

MICP4

1. Follow the instruction below and then report how the performance changed. (apply all at once) Did the performance change?

The first screenshot shows the initial setup of a Keras model for image classification. The model is defined with the following layers:

- Conv2D(32, (3, 3), input_shape=X_train.shape[1:], padding='same', activation='relu')
- Dropout(0.2)
- Conv2D(32, (3, 3), activation='relu', padding='same')
- MaxPooling2D(pool_size=(2, 2))
- Conv2D(64, (3, 3), activation='relu', padding='same')
- Dropout(0.2)
- Conv2D(64, (3, 3), activation='relu', padding='same')
- MaxPooling2D(pool_size=(2, 2))
- Conv2D(128, (3, 3), activation='relu', padding='same')
- Dropout(0.2)
- Conv2D(128, (3, 3), activation='relu', padding='same')
- MaxPooling2D(pool_size=(2, 2))
- Flatten()
- Dropout(0.2)
- Dense(1024, activation='relu', kernel_constraint=maxnorm(3))
- Dropout(0.2)
- Dense(512, activation='relu', kernel_constraint=maxnorm(3))
- Dropout(0.2)
- Dense(num_classes, activation='softmax')

The training process is initiated with the following parameters:

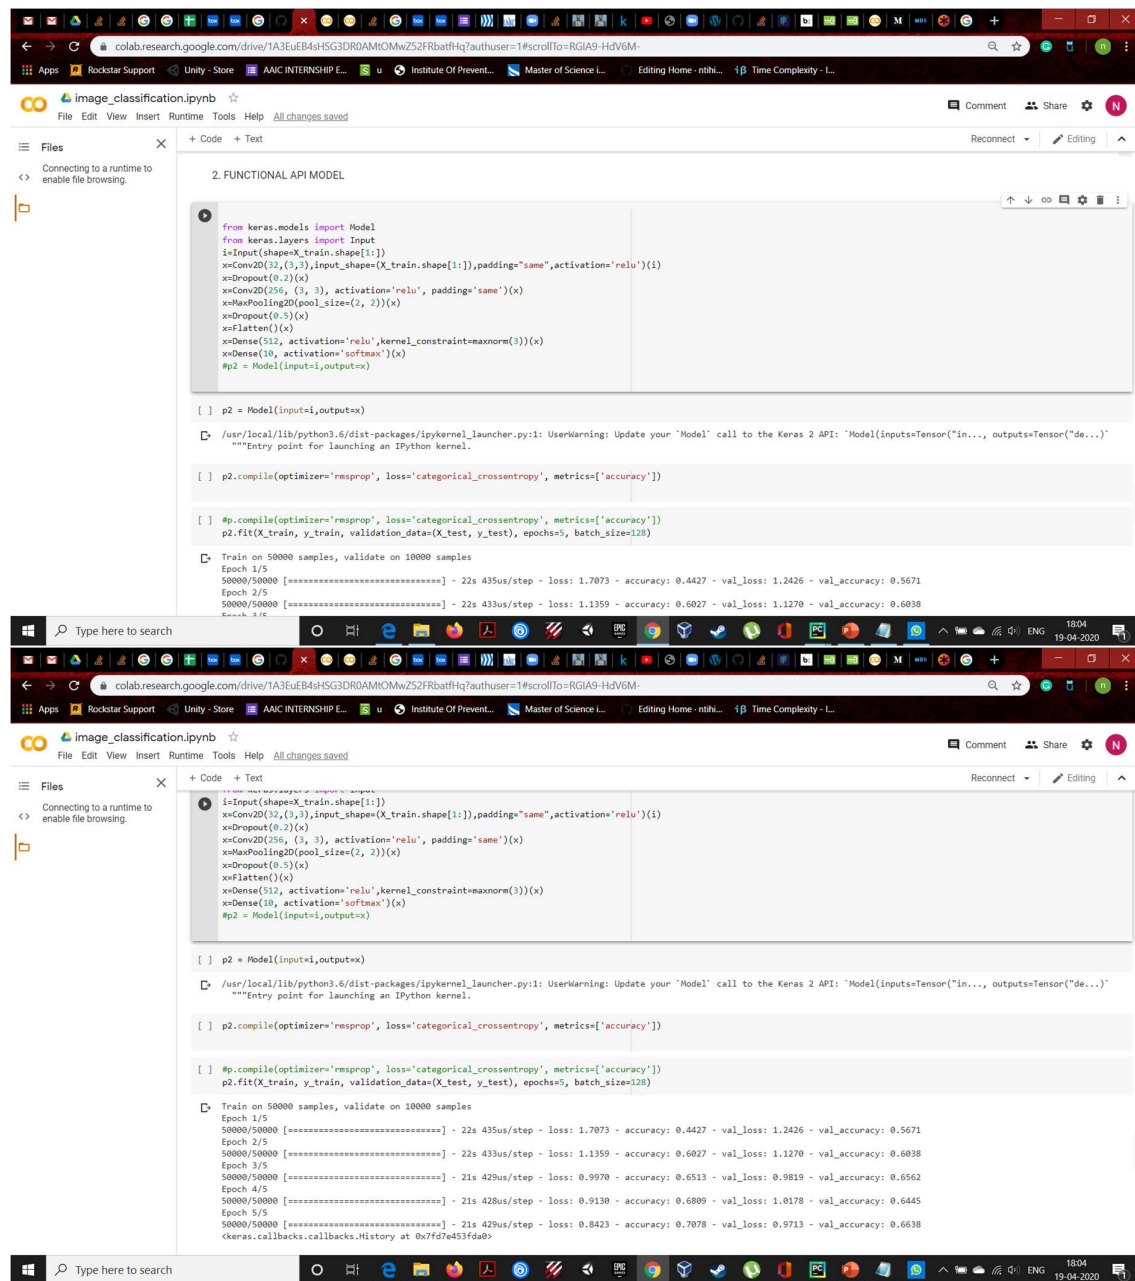
- epochs = 5
- lr_rate = 0.001
- decay = lr_rate/epochs
- sgd = Adam(lr=lr_rate)
- metrics=['accuracy']
- optimizer=sgd
- validation_data=(X_test, y_test)
- batch_size=128

The second screenshot shows the training progress over 20 epochs. The model's performance is tracked using loss and accuracy. The training loss decreases from 0.7231 to 0.3080, while the validation accuracy increases from 0.7372 to 0.7980. The model's performance is stable, indicating that it has reached a good fit.

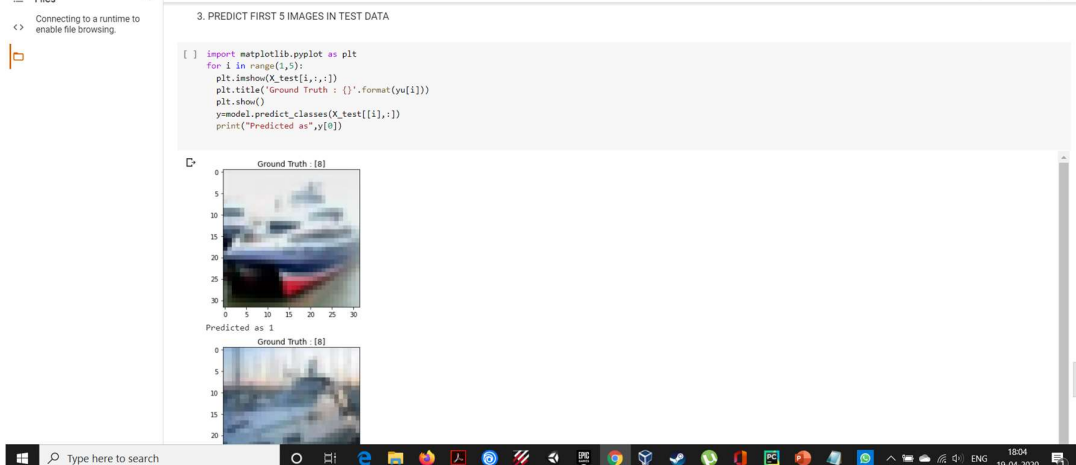

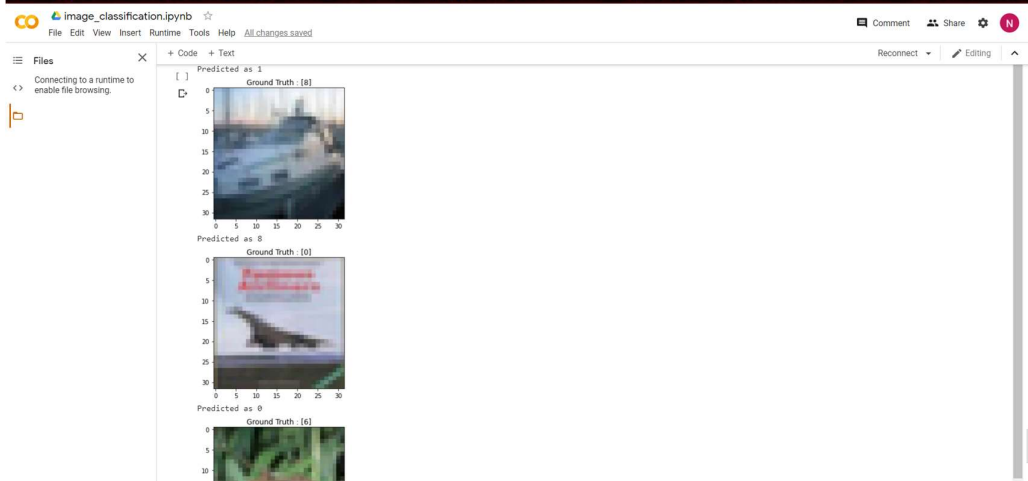
Epoch	Loss	Accuracy	Val Loss	Val Accuracy
6/20	0.7231	0.7470	0.7625	0.7372
7/20	0.6665	0.7653	0.6776	0.7663
8/20	0.6019	0.7878	0.6814	0.7687
9/20	0.5621	0.8012	0.6563	0.7767
10/20	0.5216	0.8150	0.6325	0.7821
11/20	0.4935	0.8254	0.6587	0.7792
12/20	0.4552	0.8397	0.6445	0.7933
13/20	0.4342	0.8470	0.6719	0.7814
14/20	0.4107	0.8529	0.6254	0.7980
15/20	0.3791	0.8660	0.6317	0.7976
16/20	0.3624	0.8720	0.6360	0.7952
17/20	0.3490	0.8756	0.6469	0.7951
18/20	0.3311	0.8826	0.6390	0.8017
19/20	0.3197	0.8866	0.6510	0.8046
20/20	0.3080	0.8926	0.6516	0.7980

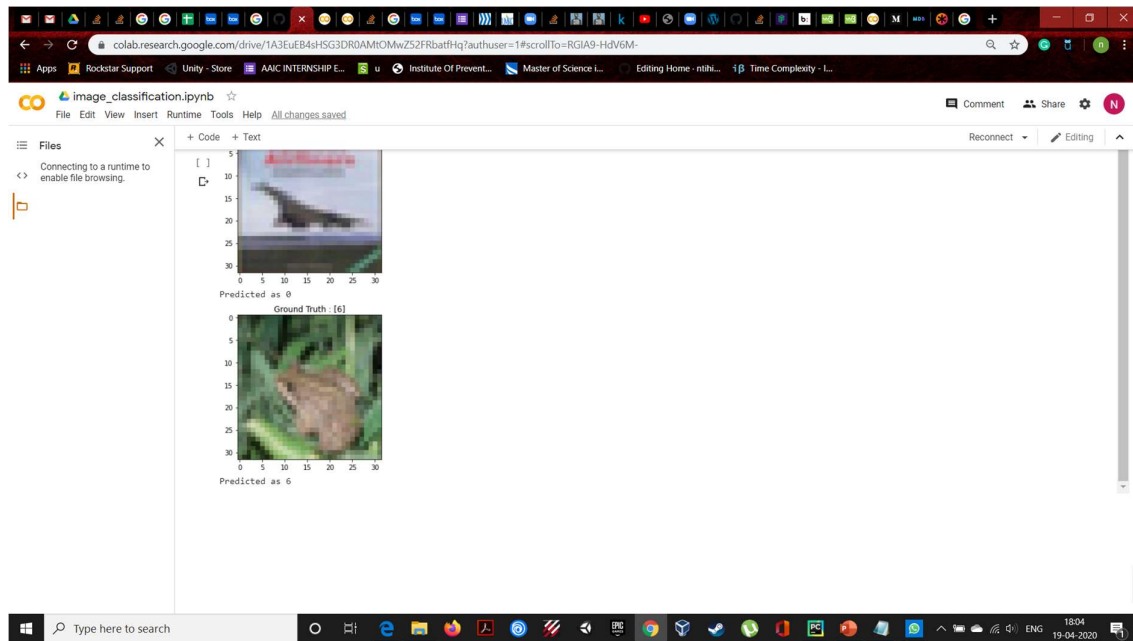
Yes, the performance has changed. The model got it best fit in minimal epochs as we can see for only 20 epochs the model got overfitted

2. Change the previous model into Keras API model.



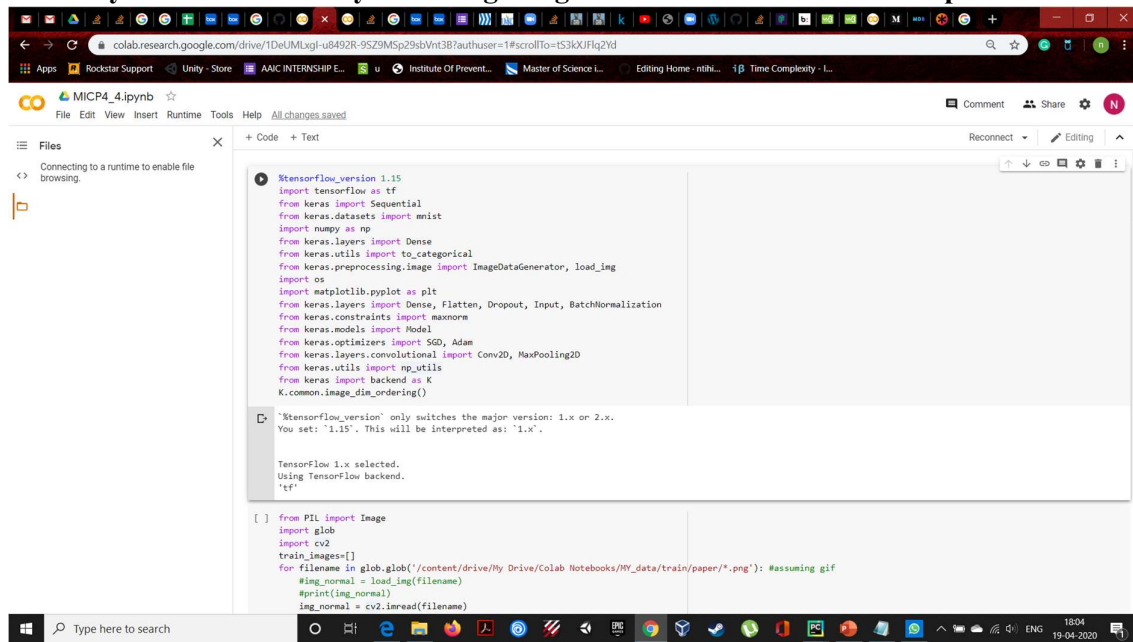
3. predict the first 4 image of the test data. Then, print the actual label for those 4 images (label means the probability associated with them) to check if the model predicted correctly or not





For the first image it has predicted wrong .

4.Build your own dataset by collecting images from the internet for example



e:

The screenshot shows a Google Colab notebook interface. The browser address bar displays a URL from colab.research.google.com. The notebook title is 'MICP4_4.ipynb'. The left sidebar shows a 'Files' panel with a message 'Connecting to a runtime to enable file browsing.' The main code area contains the following Python code:

```
[ ] import random
random.shuffle(train_images)

[ ] x_train=[]
y_train=[]
for im,label in train_images:
    x_train.append(im)
    y_train.append(label)

[ ] from PIL import Image
import glob
import cv2
test_images=[]
for filename in glob.glob('/content/drive/My Drive/Colab Notebooks/My_data/test/paper/*.png'): #assuming gif
    #img_normal = load_img(filename)
    #print(img_normal)
    img_normal = cv2.imread(filename)
    #print(img.shape)
    #print(type(img.shape))
    output = cv2.resize(img_normal, (32,32))
    test_images.append([output,0])

[ ] for filename in glob.glob('/content/drive/My Drive/Colab Notebooks/My_data/test/rock/*.png'): #assuming gif
    #img_normal = load_img(filename)
    #print(img_normal)
    img_normal = cv2.imread(filename)
    #print(img.shape)
    #print(type(img.shape))
    output = cv2.resize(img_normal, (32,32))
    test_images.append([output,1])

[ ] for filename in glob.glob('/content/drive/My Drive/Colab Notebooks/My_data/test/scissors/*.jpg'): #assuming gif
    #img_normal = load_img(filename)
    #print(img_normal)
```

The bottom of the image shows a Windows taskbar with a search bar and various application icons. The system clock indicates 18:04 on 19-04-2020.




```
from PIL import Image
import glob
import cv2
train_images=[]
for filename in glob.glob('/content/drive/My Drive/Colab Notebooks/My_data/train/paper/*.png'): #assuming gif
    img_normal = load_img(filename)
    #print(img_normal)
    img_normal = cv2.imread(filename)
    #print(img.shape)
    #print(type(img.shape))
    output = cv2.resize(img_normal, (32,32))
    train_images.append([output,0])

for filename in glob.glob('/content/drive/My Drive/Colab Notebooks/My_data/train/rock/*.png'): #assuming gif
    img_normal = load_img(filename)
    #print(img_normal)
    img_normal = cv2.imread(filename)
    #print(img.shape)
    #print(type(img.shape))
    output = cv2.resize(img_normal, (32,32))
    train_images.append([output,1])

for filename in glob.glob('/content/drive/My Drive/Colab Notebooks/My_data/train/scissors/*.jpg'): #assuming gif
    img_normal = load_img(filename)
    #print(img_normal)
    img_normal = cv2.imread(filename)
    #print(img.shape)
    #print(type(img.shape))
    output = cv2.resize(img_normal, (32,32))
    train_images.append([output,2])

import random
random.shuffle(train_images)
```



```
x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train = x_train / 255.0
x_test = x_test / 255.0

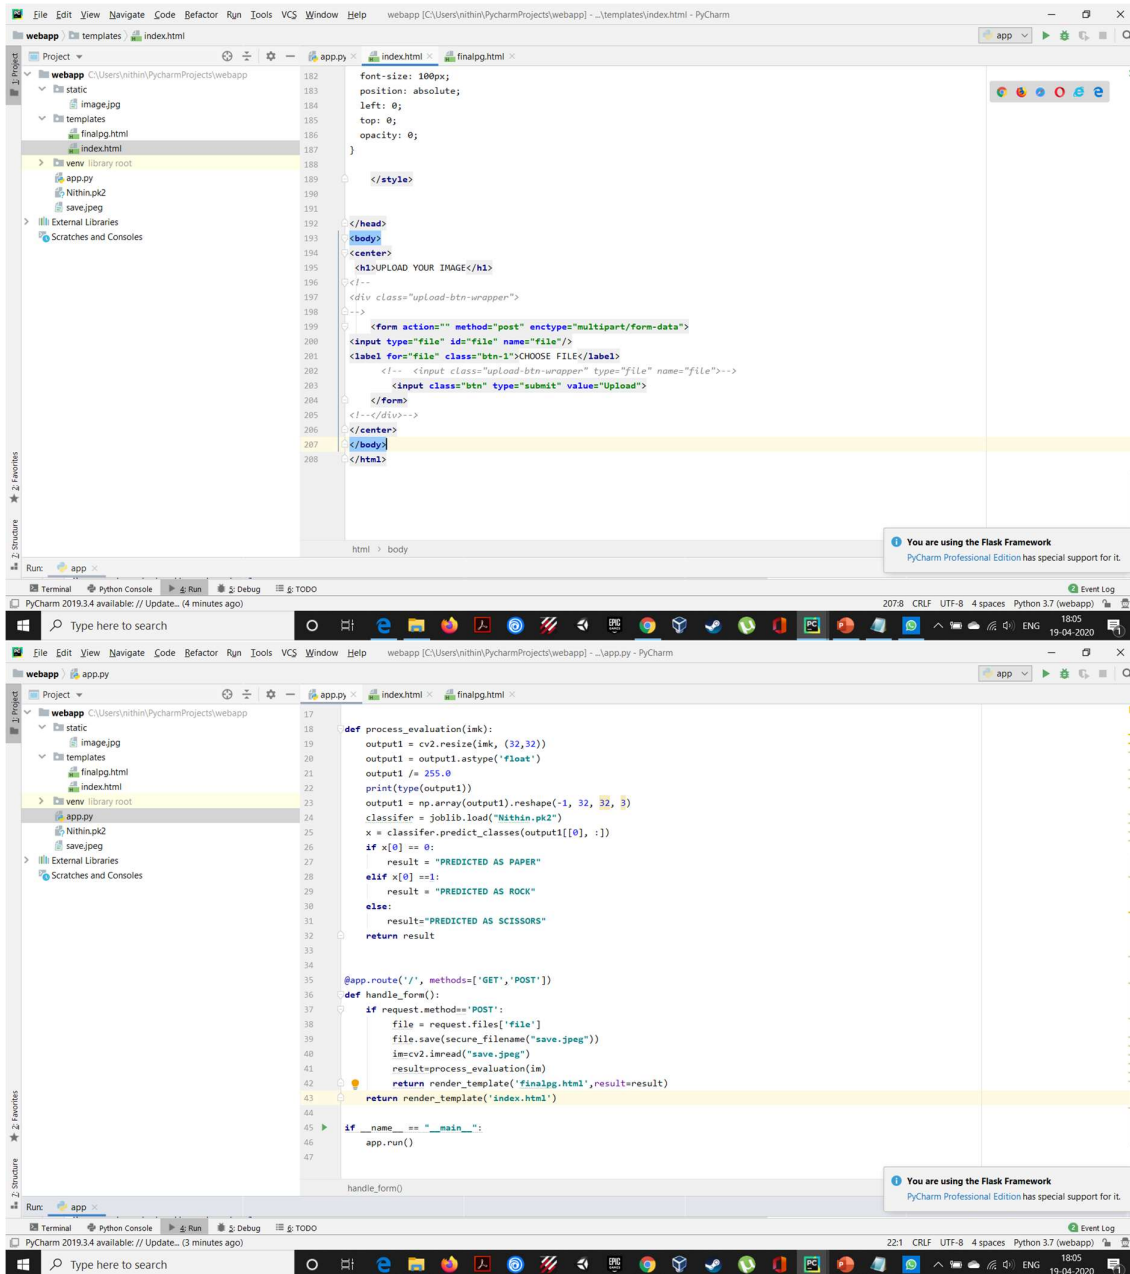
y_train = np_utils.to_categorical(y_train)
y_test = np_utils.to_categorical(y_test)
num_classes = y_test.shape[1]

model = Sequential()
model.add(Conv2D(256, (3, 3), input_shape=(x_train.shape[1:]), padding='same', activation='relu'))
model.add(Dropout(0.5))
model.add(Conv2D(256, (3, 3), activation='relu', padding='same'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.5))
model.add(Flatten())
model.add(Dense(512, activation='relu', kernel_constraint=MaxNorm(3)))
model.add(Dropout(0.3))
model.add(Dense(num_classes, activation='softmax'))

WARNING:tensorflow:From /tensorflow-1.15.2/python3.6/tensorflow_core/python/ops/resource_variable_ops.py:1630: calling BaseResourceVariable.__init__ (from tensorflow.p
Instructions for updating:
If using Keras pass *_constraint arguments to layers.
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:4070: The name tf.nn.max_pool is deprecated. Please use tf.nn.max_po

epochs = 20
lr_rate = 0.001
decay = lr_rate/epochs
sgd = Adam(lr=lr_rate)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])

model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=epochs, batch_size=128)
```

```
[ ] epochs = 20
lr_rate = 0.001
decay = lr_rate/epoch
sgd = Adam(lr=lr_rate)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])

[ ] model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=epochs, batch_size=128)

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:422: The name tf.global_variables is deprecated. Please use tf.compat_v1.get_global_variables instead.

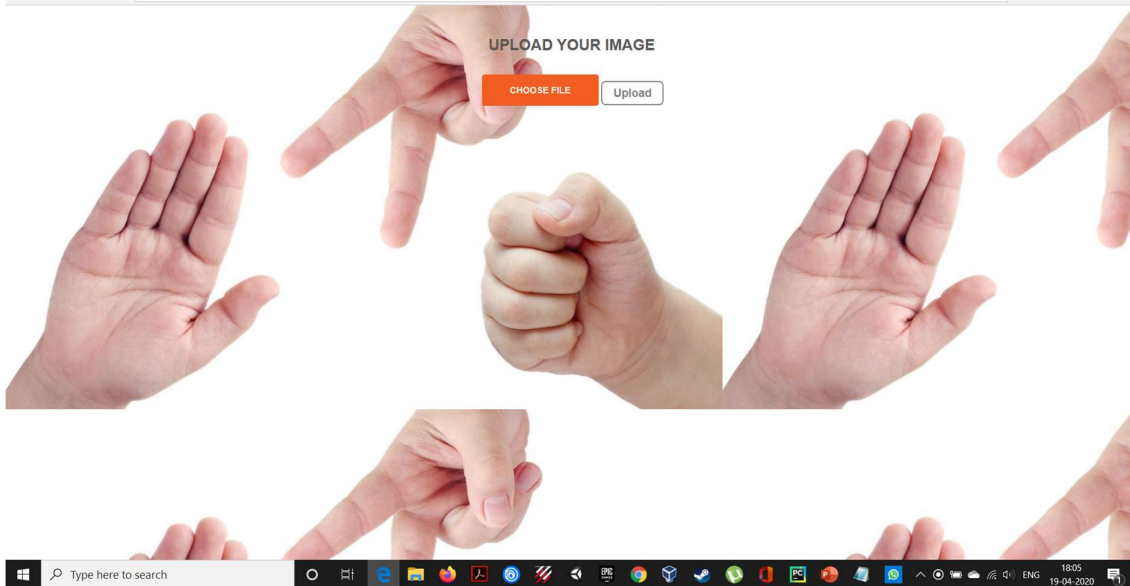
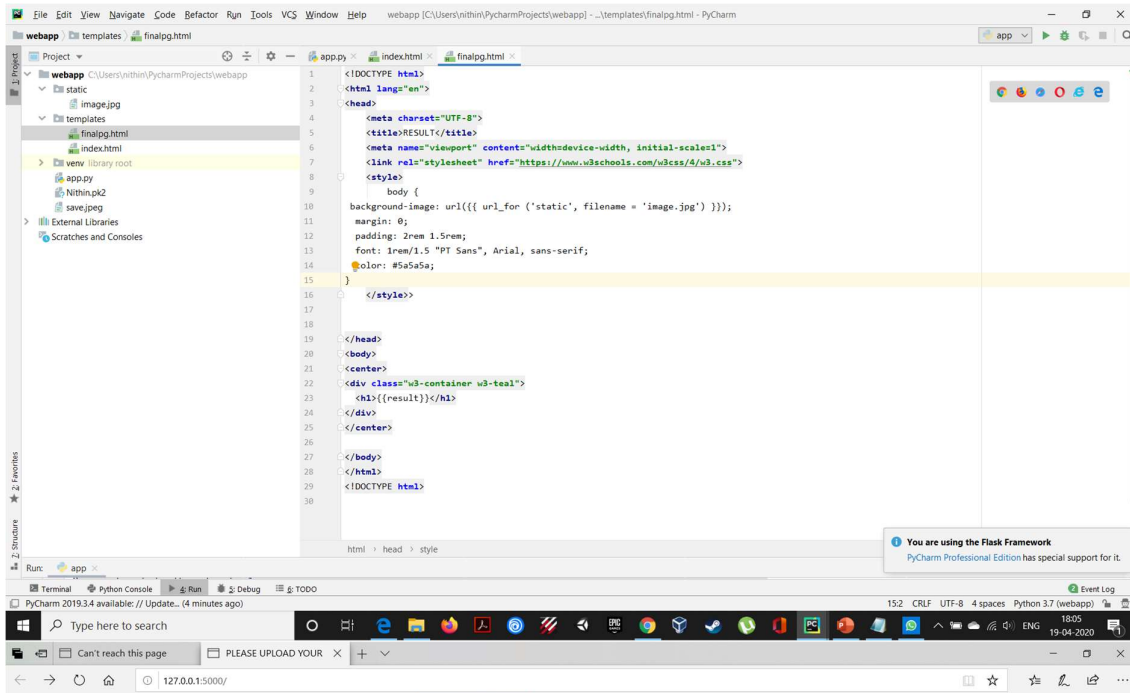
Train on 1344 samples, validate on 1344 samples
Epoch 1/20
1344/1344 [=====] - 140s 104ms/step - loss: 1.6107 - accuracy: 0.4888 - val_loss: 0.6930 - val_accuracy: 0.5119
Epoch 2/20
1344/1344 [=====] - 135s 101ms/step - loss: 0.6911 - accuracy: 0.5171 - val_loss: 0.6845 - val_accuracy: 0.5119
Epoch 3/20
1344/1344 [=====] - 135s 101ms/step - loss: 0.6694 - accuracy: 0.5982 - val_loss: 0.6074 - val_accuracy: 0.8125
Epoch 4/20
1344/1344 [=====] - 134s 100ms/step - loss: 0.5184 - accuracy: 0.7872 - val_loss: 0.4005 - val_accuracy: 0.8438
Epoch 5/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.4367 - accuracy: 0.8095 - val_loss: 0.3645 - val_accuracy: 0.8757
Epoch 6/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.3820 - accuracy: 0.8430 - val_loss: 0.2916 - val_accuracy: 0.8936
Epoch 7/20
1344/1344 [=====] - 134s 100ms/step - loss: 0.3345 - accuracy: 0.8772 - val_loss: 0.2585 - val_accuracy: 0.9055
Epoch 8/20
1344/1344 [=====] - 134s 100ms/step - loss: 0.3328 - accuracy: 0.8728 - val_loss: 0.2660 - val_accuracy: 0.9152
Epoch 9/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.3139 - accuracy: 0.8631 - val_loss: 0.2755 - val_accuracy: 0.9129
Epoch 10/20
1344/1344 [=====] - 139s 104ms/step - loss: 0.3123 - accuracy: 0.8772 - val_loss: 0.2476 - val_accuracy: 0.9263
Epoch 11/20
1344/1344 [=====] - 136s 101ms/step - loss: 0.3088 - accuracy: 0.8765 - val_loss: 0.2181 - val_accuracy: 0.9226
Epoch 12/20
1344/1344 [=====] - 137s 102ms/step - loss: 0.2973 - accuracy: 0.8869 - val_loss: 0.1893 - val_accuracy: 0.9308
Epoch 13/20
1344/1344 [=====] - 138s 103ms/step - loss: 0.2752 - accuracy: 0.8891 - val_loss: 0.2205 - val_accuracy: 0.9397
Epoch 14/20
1344/1344 [=====] - 141s 105ms/step - loss: 0.2736 - accuracy: 0.8943 - val_loss: 0.1932 - val_accuracy: 0.9427
Epoch 15/20
1344/1344 [=====] - 137s 102ms/step - loss: 0.2476 - accuracy: 0.9040 - val_loss: 0.1714 - val_accuracy: 0.9360
Epoch 16/20
1344/1344 [=====] - 136s 102ms/step - loss: 0.2671 - accuracy: 0.8966 - val_loss: 0.1607 - val_accuracy: 0.9509
Epoch 17/20
1344/1344 [=====] - 137s 102ms/step - loss: 0.2445 - accuracy: 0.8981 - val_loss: 0.1666 - val_accuracy: 0.9568
Epoch 18/20
1344/1344 [=====] - 139s 103ms/step - loss: 0.2359 - accuracy: 0.9025 - val_loss: 0.1831 - val_accuracy: 0.9516
Epoch 19/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.2120 - accuracy: 0.9219 - val_loss: 0.1401 - val_accuracy: 0.9516
Epoch 20/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.2188 - accuracy: 0.9144 - val_loss: 0.1375 - val_accuracy: 0.9568
keras.callbacks.callbacks.History at 0x7fed2cf9e358>
```

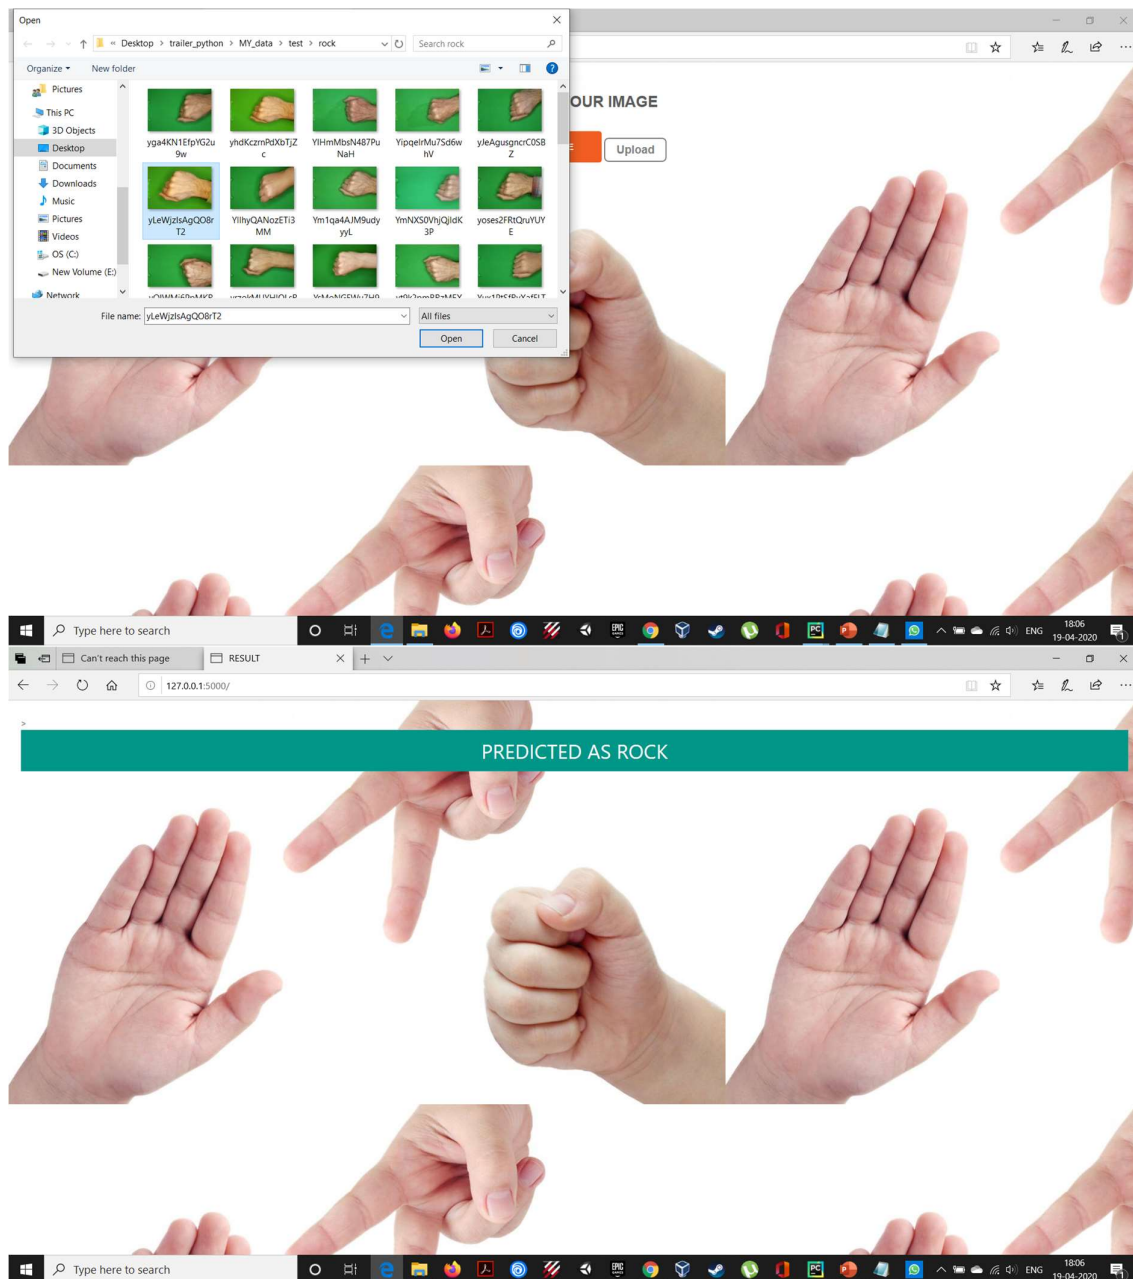
```
[ ] epochs = 20
lr_rate = 0.001
decay = lr_rate/epoch
sgd = Adam(lr=lr_rate)
model.compile(loss='categorical_crossentropy', optimizer=sgd, metrics=['accuracy'])

[ ] model.fit(x_train, y_train, validation_data=(x_test, y_test), epochs=epochs, batch_size=128)

WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/keras/backend/tensorflow_backend.py:422: The name tf.global_variables is deprecated. Please use tf.compat_v1.get_global_variables instead.

Train on 1344 samples, validate on 1344 samples
Epoch 1/20
1344/1344 [=====] - 140s 104ms/step - loss: 1.6107 - accuracy: 0.4888 - val_loss: 0.6930 - val_accuracy: 0.5119
Epoch 2/20
1344/1344 [=====] - 135s 101ms/step - loss: 0.6911 - accuracy: 0.5171 - val_loss: 0.6845 - val_accuracy: 0.5119
Epoch 3/20
1344/1344 [=====] - 135s 101ms/step - loss: 0.6694 - accuracy: 0.5982 - val_loss: 0.6074 - val_accuracy: 0.8125
Epoch 4/20
1344/1344 [=====] - 134s 100ms/step - loss: 0.5184 - accuracy: 0.7872 - val_loss: 0.4005 - val_accuracy: 0.8438
Epoch 5/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.4367 - accuracy: 0.8095 - val_loss: 0.3645 - val_accuracy: 0.8757
Epoch 6/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.3820 - accuracy: 0.8430 - val_loss: 0.2916 - val_accuracy: 0.8936
Epoch 7/20
1344/1344 [=====] - 134s 100ms/step - loss: 0.3345 - accuracy: 0.8772 - val_loss: 0.2585 - val_accuracy: 0.9055
Epoch 8/20
1344/1344 [=====] - 134s 100ms/step - loss: 0.3328 - accuracy: 0.8728 - val_loss: 0.2660 - val_accuracy: 0.9152
Epoch 9/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.3139 - accuracy: 0.8631 - val_loss: 0.2755 - val_accuracy: 0.9129
Epoch 10/20
1344/1344 [=====] - 139s 104ms/step - loss: 0.3123 - accuracy: 0.8772 - val_loss: 0.2476 - val_accuracy: 0.9263
Epoch 11/20
1344/1344 [=====] - 136s 101ms/step - loss: 0.3088 - accuracy: 0.8765 - val_loss: 0.2181 - val_accuracy: 0.9226
Epoch 12/20
1344/1344 [=====] - 137s 102ms/step - loss: 0.2973 - accuracy: 0.8869 - val_loss: 0.1893 - val_accuracy: 0.9308
Epoch 13/20
1344/1344 [=====] - 138s 103ms/step - loss: 0.2752 - accuracy: 0.8891 - val_loss: 0.2205 - val_accuracy: 0.9397
Epoch 14/20
1344/1344 [=====] - 141s 105ms/step - loss: 0.2736 - accuracy: 0.8943 - val_loss: 0.1932 - val_accuracy: 0.9427
Epoch 15/20
1344/1344 [=====] - 137s 102ms/step - loss: 0.2476 - accuracy: 0.9040 - val_loss: 0.1714 - val_accuracy: 0.9360
Epoch 16/20
1344/1344 [=====] - 136s 102ms/step - loss: 0.2671 - accuracy: 0.8966 - val_loss: 0.1607 - val_accuracy: 0.9509
Epoch 17/20
1344/1344 [=====] - 137s 102ms/step - loss: 0.2445 - accuracy: 0.8981 - val_loss: 0.1666 - val_accuracy: 0.9568
Epoch 18/20
1344/1344 [=====] - 139s 103ms/step - loss: 0.2359 - accuracy: 0.9025 - val_loss: 0.1831 - val_accuracy: 0.9516
Epoch 19/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.2120 - accuracy: 0.9219 - val_loss: 0.1401 - val_accuracy: 0.9516
Epoch 20/20
1344/1344 [=====] - 135s 100ms/step - loss: 0.2188 - accuracy: 0.9144 - val_loss: 0.1375 - val_accuracy: 0.9568
keras.callbacks.callbacks.History at 0x7fed2cf9e358>
```



BY,
DUKKIPATI SRI SAI NITHIN CHOWDARY
CLASS ID: 04